

OCULUS VR, INC

SDK Overview
Document Version 1.2

Authors:

Michael Antonov
Nate Mitchell
Andrew Reisse
Lee Cooper
Steve LaValle

Date:

April 24, 2013

2013 Oculus VR, Inc. All rights reserved.

Oculus VR, Inc. 19800 MacArthur Blvd Suite 450 Irvine, CA 92612

Except as otherwise permitted by Oculus VR, Inc., this publication, or parts thereof, may not be reproduced in any form, by any method, for any purpose. Certain materials included in this publication are reprinted with the permission of the copyright holder.

All brand names, product names or trademarks belong to their respective holders.

Disclaimer

THIS PUBLICATION AND THE INFORMATION CONTAINED HEREIN IS MADE AVAILABLE BY OCULUS VR, INC. AS IS. OCULUS VR, INC. DISCLAIMS ALL WARRANTIES, EITHER EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE REGARDING THESE MATERIALS.

Contents

1	Introduction	5
2	Oculus Rift Hardware Setup	5
2.1	Display Specifications	5
2.2	Tracker Specifications	5
2.3	Additional Vision Lenses	6
2.3.1	Changing vision lenses	6
2.4	Screen Distance Adjustment	6
2.4.1	Adjusting the screen distance	7
2.5	Control Box Setup	7
2.5.1	Adjusting brightness and contrast	7
2.6	Monitor Setup	8
3	Oculus Rift SDK Setup	9
3.1	System Requirements	9
3.1.1	Operating systems	9
3.1.2	Minimum system requirements	9
3.2	Installation	10
3.3	Directory Structure	10
3.4	Compiler Settings	10
3.5	Makefiles, Projects, and Build Solutions	10
3.5.1	Windows	10
3.5.2	MacOS	11
3.6	Terminology	11
4	Getting Started	12
4.1	OculusWorldDemo	12
4.1.1	Controls	12
4.1.2	Using OculusWorldDemo	13
4.2	Using the SDK Beyond the OculusWorldDemo	14
4.2.1	Software developers and integration engineers	14

4.2.2	Artists and game designers	14
5	LibOVR Integration Tutorial	16
5.1	Outline of Integration Tasks	16
5.2	Initialization of LibOVR	16
5.3	Sensor Data and Head Tracking	18
5.3.1	Some details on sensor fusion	20
5.3.2	Magnetometer-based yaw error correction	21
5.4	User Input Integration	23
5.5	Rendering Configuration	23
5.5.1	Rendering stereo	24
5.5.2	Distortion correction	28
5.5.3	Distortion scale	30
5.5.4	Distortion and FOV	31
5.5.5	StereoConfig utility class	32
5.5.6	Rendering performance	33
6	Optimization	34
6.1	Latency	34
A	Display Device Management	36
A.1	Display Identification	36
A.2	Display Configuration	36
A.2.1	Duplicate display mode	37
A.2.2	Extended display mode	37
A.2.3	Standalone display mode	37
A.3	Selecting A Display Device	37
A.3.1	Windows	37
A.3.2	MacOS	39
A.4	Rift Display Considerations	40
A.4.1	Duplicate mode VSync	40
A.4.2	Extended mode problems	40
A.4.3	Observing Rift output on a monitor	40

A.4.4	Windows: Direct3D enumeration	41
B	Chromatic Aberration	42
B.1	Correction	42
B.2	Sub-channel aberration	42
B.3	Shader implementation	42

1 Introduction

Thanks for downloading the Oculus Software Development Kit (SDK)!

This document will detail how to install, configure, and use the Oculus SDK. The Oculus SDK includes all of the components that developers need to integrate the Oculus Rift with their game engine or application. The core of the SDK is made up of source code and binary libraries. The Oculus SDK also includes documentation, samples, and tools to help developers get started.

This document focuses on the C++ API of the Oculus SDK. Integration with the Unreal Engine 3 (UE3) and Unity game engine is available as follows:

- Unity integration is available as a separate package from the [Oculus Developer Center](#).
- Unreal Engine 3 integration is also available as a separate package from the Oculus Developer Center. You will need a full UE3 license to access the version of Unreal with Oculus integration. If you have a full UE3 license, you can email support@oculusvr.com to be granted download access.

2 Oculus Rift Hardware Setup

In addition to the Oculus SDK, you will also need the hardware provided by the Oculus Rift Development Kit (DK). The DK includes an Oculus Rift development headset (Rift), control box, required cabling, and additional pairs of lenses for different vision characteristics.

2.1 Display Specifications

- 7 inch diagonal viewing area
- 1280 × 800 resolution (720p). This is split between both eyes, yielding 640 × 800 per eye.
- 64mm fixed distance between lens centers
- 60Hz LCD panel
- DVI-D Single Link
- HDMI 1.3+
- USB 2.0 Full Speed+

2.2 Tracker Specifications

- Up to 1000Hz sampling rate
- Three-axis gyroscope, which senses angular velocity
- Three-axis magnetometer, which senses magnetic fields
- Three-axis accelerometer, which senses accelerations, including gravitational

2.3 Additional Vision Lenses

The Rift comes installed with lenses for users with 20/20 or farsighted vision. If your vision is 20/20 or farsighted, you won't need to change your lenses and you can proceed to Section 2.4.

For nearsighted users, two additional pairs of lenses are included with the kit. Although they may not work perfectly for all nearsighted users, they should enable most people to use the headset without glasses or contact lenses.

The medium-depth lenses are for users who are moderately nearsighted. The shortest-depth lenses are for users who are very nearsighted. We recommend that users experiment with the different lenses to find the ones that work best for them.

The lenses are also marked with the letters 'A', 'B', and 'C' to aid identification. The recommended lenses are as follows:

Lenses	Appropriate Vision
A	20/20 or farsighted
B	Moderately nearsighted
C	Very nearsighted

Note: If your eyes have special characteristics such as astigmatism, the provided lenses may not be sufficient to correct your vision. In this case, we recommend wearing contact lenses or glasses. Note, however, that using glasses will cut down on your effective field of view.

2.3.1 Changing vision lenses

Note: Changing the lens may cause dust or debris to get inside the Rift. We strongly recommend changing the lenses in the cleanest space possible! Do not store the Rift without lenses installed.

To change lenses, first turn the headset upside down (this is to minimize the amount of dust and debris that can enter the headset) and gently unscrew the lenses currently attached to the headset. Unscrewing the lenses doesn't require much pressure; a light touch is most effective. The right lens unscrews clockwise. The left lens unscrews counterclockwise.

Place the old lenses in a safe place, then take the new lenses and install them the same way you removed the original pair. Remember to keep your headset upside down during this process. Once the new lenses are securely in place, you're all set!

After changing the lenses, you may need to adjust the distance of the assembly that holds the screen and lenses closer or farther away from your face. This is covered next.

2.4 Screen Distance Adjustment

The headset has an adjustment feature that allows you to change the distance of the fixture that holds the screen and lenses from your eyes. This is provided to accommodate different facial characteristics and vision lenses. For example, if the lenses are too close to your eyes, then you should adjust the fixture outward, moving the lenses and the screen away from your face. You can also use this to provide more room for eyeglasses.

Note: Everyone should take some time to adjust the headset for maximum comfort. While doing so, an important consideration is that the lenses should be situated as close to your eyes as possible. Remember that the maximal field of view occurs when your eyes are as close to the lenses as possible without actually touching them.

2.4.1 Adjusting the screen distance

There are two screw mechanisms of either side of the headset that can be adjusted using a coin. These screws control the location of the screen assembly. The setting for the two screw mechanisms should always match unless you're in the process of adjusting them.

Turn the screw mechanism toward the lenses to bring the assembly closer to the user. Turn the screw mechanism toward the display to move the assembly farther away from the user. After changing one side, ensure that the other side is turned to the same setting!

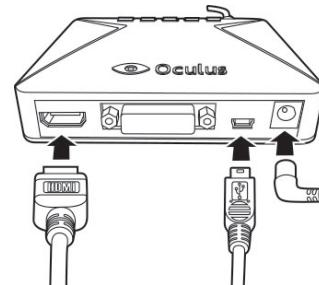
2.5 Control Box Setup

The headset is connected to the control box by a 6ft cable. The control box takes in video, USB, and power, and sends them out over a single cord to minimize the amount of cabling running to the headset.

1. Connect one end of the video cable (DVI or HDMI) to your computer and the other end to the control box.

Note: There should only be one video-out cable running to the control box at a time (DVI or HDMI, not both).

2. Connect one end of the USB cable to your computer and the other to the control box.
3. Plug the power cord into an outlet and connect the other end to the control box.



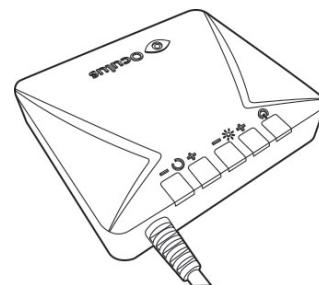
You can power on the DK using the power button on the top of the control box. A blue LED indicates whether the DK is powered on or off. The Rift screen will only stay on when all three cables are connected.

2.5.1 Adjusting brightness and contrast

The brightness and contrast of the headset can be adjusted using the buttons on the top of the control box.

Looking from the back side:

- The leftmost buttons adjust the display's contrast.
- The neighboring two adjust the display's brightness.
- The rightmost button turns the power on and off.



2.6 Monitor Setup

Once the Oculus Rift is connected to your computer, it should be automatically recognized as an additional monitor and Human Input Device (HID).

The Rift can be set to mirror or extend your current monitor setup using your computer's display settings. We recommend using the Rift as an extended monitor in most cases, but it's up to you to decide which configuration works best for you. This is covered in more detail in Appendix A. Regardless of the monitor configuration, is it currently not possible to see the desktop clearly inside the Rift. This would require stereo rendering and distortion correction, which is only available while rendering the game scene.

Whether you decide to mirror or extend your desktop, the resolution of the Rift should always be set to 1280×800 (720p).

3 Oculus Rift SDK Setup

3.1 System Requirements

3.1.1 Operating systems

The Oculus SDK currently supports MacOS, Windows Vista, Windows 7, and Windows 8.

3.1.2 Minimum system requirements

There are no specific computer hardware requirements for the Oculus SDK; however, we recommend that developers use a computer with a modern graphics card. A good benchmark is to try running Unreal Engine 3 and Unity at 60 frames per second (FPS) with vertical sync and stereo 3D enabled. If this is possible without dropping frames, then your configuration should be sufficient for Oculus Rift development!

The following components are provided as a guideline:

- Windows: Vista, 7, or 8
- MacOS: 10.6+
- 2.0+ GHz processor
- 2 GB system RAM
- Direct3D10 or OpenGL 3 compatible video card.

Although many lower end and mobile video cards, such as the Intel HD 4000, have the shader and graphics capabilities to run minimal Rift demos, their rendering throughput may be inadequate for full-scene 60 FPS VR rendering with stereo and distortion. Developers targeting this hardware will need to be very conscious of scene geometry because low-latency rendering at 60 FPS is critical for a usable VR experience.

If you are looking for a portable VR workstation, we've found that the Nvidia 650M inside of a MacBook Pro Retina provides enough graphics power for our demo development.

3.2 Installation

The latest version of the Oculus SDK is available at <http://developer.oculusvr.com>.

The naming convention for the Oculus SDK release package is `ovr_packagetype_major.minor.build`. For example, the initial build was `ovr_lib_0.1.1.zip`.

3.3 Directory Structure

The installed Oculus SDK package contains the following subdirectories:

3rdParty	Third party SDK components used by samples, such as TinyXml.
DOC	SDK Documentation, including this document.
LibOVR	Libraries, source code, projects, and makefiles for the SDK.
LibOVR/Include	Public include header files, including OVR.h. Header files here reference other headers in LibOVR/Src.
LibOVR/Lib	Pre-built libraries for use in your project.
LibOVR/Src	Source code and internally referenced headers.
Samples	Samples that integrate and leverage the Oculus SDK.

3.4 Compiler Settings

The LibOVR libraries do not require exception handling or RTTI support, thereby allowing your game to disable these features for efficiency.

3.5 Makefiles, Projects, and Build Solutions

Developers can rebuild the samples and LibOVR using the projects and solutions in the Samples and LibOVR/Projects directory.

3.5.1 Windows

The Visual Studio 2010 solution and project files are provided with the SDK:

- Samples/LibOVR_Samples_Msvc2010.sln is the main solution that allows you to build and run all of the samples.
- LibOVR/Projects/Win32 contains the project needed to build the LibOVR library itself (for developers that have access to the full source).

3.5.2 MacOS

The included Xcode project `Samples/LibOVR_With_Samples.xcodeproj` allows you to build and run all of the samples, and libovr itself. The project is setup to build universal binaries (x86 and x86_64) for all recent MacOS versions (10.6 and newer).

3.6 Terminology

The following terms are crucial in the coming sections:

Interpupillary distance (IPD)	The distance between the eye pupils. The default value in the SDK is 64mm which corresponds to the average human distance, but values of 54mm to 72mm are possible.
Field of view (FOV)	The full vertical viewing angle used to configure rendering. This is computed based on the eye distance and display size.
Aspect ratio	The ratio of horizontal resolution to vertical resolution. The aspect ratio for each eye on the Oculus Rift is 0.8.
Multisampling	Hardware anti-aliasing mode supported by many video cards.

4 Getting Started

Your developer kit is unpacked and plugged in, you've installed the SDK, and you are ready to go. Where is the best place to begin?

If you haven't already, take a moment to adjust the Rift headset so that it's comfortable for your head and eyes. More detailed information about configuring the Rift can be found in Section 2.

Once your hardware is fully configured, the next step is to test the development kit. The SDK comes with a set of full-source C++ samples designed to help developers get started quickly. These include:

- **OculusWorldDemo** - A visually appealing Tuscany scene with on-screen text and controls.
- **OculusRoomTiny** - A minimal C++ sample showing sensor integration and rendering on the Rift.
- **SensorBoxTest** - A 3D rendered box that demonstrates sensor fusion by tracking and displaying the rotation of the Rift.

We recommend running the pre-built OculusWorldDemo as a first-step in exploring the SDK. You can find a link to the Windows executable in the root of the Oculus SDK installation.

4.1 OculusWorldDemo

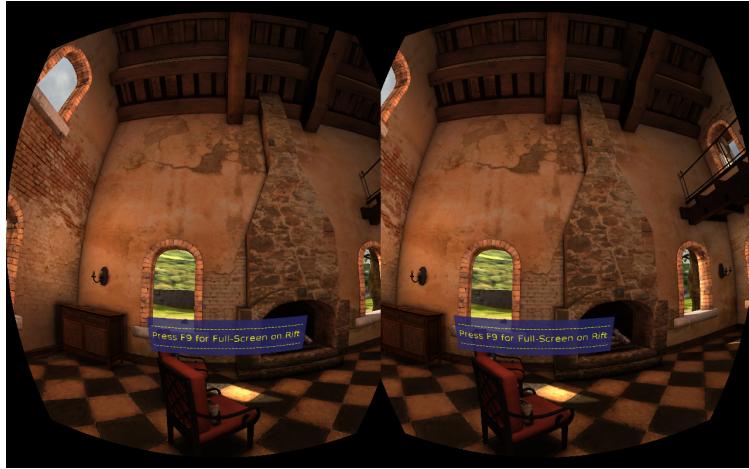


Figure 1: Screenshot of the OculusWorldDemo application.

4.1.1 Controls

Key or Input	Movement	Key	Function
W, S	Move forward, back	F1	No stereo, no distortion
A, D	Strafe left, right	F2	Stereo, no distortion
Mouse move	Look left, right	F3	Stereo and distortion
Left gamepad stick	Move	F9	Hardware full-screen (low latency)
Right gamepad stick	Turn	F11	Windowed full-screen (no blinking)

Key(s)	Function	Key(s)	Function
R	Reset sensor orientation	Insert, Delete	Adjust interpupillary distance
G	Toggle grid overlay	PageUp, PageDown	Adjust aspect ratio
Spacebar	Toggle debug info overlay	[,]	Adjust field of view
Esc	Cancel full-screen	Y, H	Adjust k_1 coefficient
- , +	Adjust eye height	U, J	Adjust k_2 coefficient
Z	Start manual mag calibration	P	Toggle motion prediction
X	Start auto mag calibration	N, M	Adjust motion prediction
F6	Set mag ref and show yaw marks	C	Toggle chromatic aberration

4.1.2 Using OculusWorldDemo

Once you've launched OculusWorldDemo, take a moment to look around using the Rift and double check that all of the hardware is working properly. You should see an image similar to the screenshot in Figure 1. Press F9 or F11 to switch rendering to the Oculus Rift.

- **F9** - Switches to hardware full-screen mode. This will give best possible latency, but will blink monitors as Windows changes display settings. If no image shows up in the Rift, then press F9 again to cycle to the next monitor.
- **F11** - Instantly switches the rendering window to the Rift portion of the desktop. This mode has higher latency and no vsync, but is convenient for development.

If you're having problems (for example no image in the headset, no head tracking, and so on), then see the developer forums on the Oculus Developer Center. These should help for resolving common issues.

There are a number of interesting things to take note of during your first trip inside OculusWorldDemo. First, the level is designed “to scale”. Thus, everything appears to be roughly the same height as it would be in the real world. The sizes for everything, including the chairs, tables, doors, and ceiling, are based on measurements from real world objects. All of the units are measured in meters.

Depending on your actual height, you may feel shorter or taller than normal. The default eye-height of the player in OculusWorldDemo is 1.78 meters (5ft 10in), but this can be adjusted using the ‘+’ and ‘-’ keys.

As you may have already concluded, the scale of the world and the player is critical to an immersive VR experience. This means that players should be a realistic height, and that art assets should be sized proportionally. More details on scale can be found in the “Oculus Best Practices Guide” document. Among other things, the demo includes simulation of a basic head model, which causes head rotation to introduce additional displacement proportional to the offset of eyes from the base of the neck. This displacement is important for improving realism and reducing disorientation.

4.2 Using the SDK Beyond the OculusWorldDemo

4.2.1 Software developers and integration engineers

If you’re integrating the Oculus SDK into your game engine, we recommend starting by opening the sample projects (`Samples/LibOVR_With_Samples_Msvc2010.sln` or `Samples/LibOVR_With_Samples.xcodeproj`), building the projects, and experimenting with the provided sample code.

OculusRoomTiny is a good place to start because its source code compactly combines all critical features of the Oculus SDK. It contains logic necessary to initialize LibOVR core, access Oculus devices, implement head-tracking, sensor fusion, head modeling, stereoscopic 3D rendering, and distortion shaders.

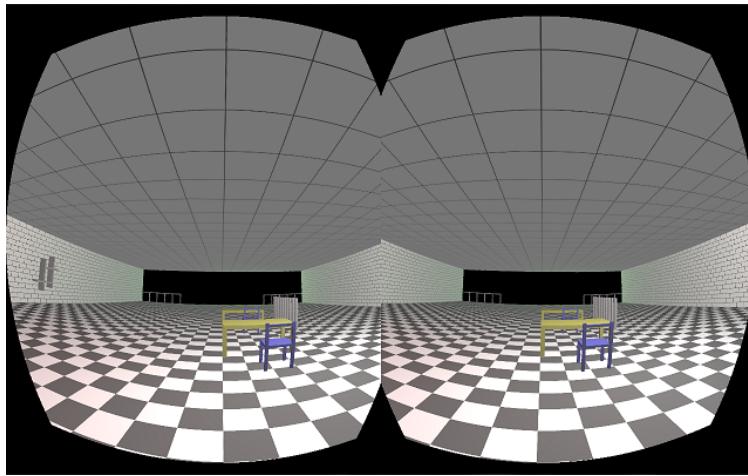


Figure 2: Screenshot of the OculusRoomTiny application.

OculusWorldDemo is a more complex sample. It is intended to be portable and supports many more features including: windowed/full-screen mode switching, XML 3D model and texture loading, movement collision detection, adjustable distortion and view key controls, 2D UI text overlays, and so on. This is a good application to experiment with once you are familiar with Oculus SDK basics.

Beyond experimenting with the provided sample code, you should continue to follow this document. We’ll cover important topics including the Oculus kernel, initializing devices, head-tracking, rendering for the Rift, and minimizing latency.

4.2.2 Artists and game designers

If you’re an artist or game designer unfamiliar in C++, we recommend downloading UE3 or Unity along with the corresponding Oculus integration. You can use our out-of-the-box integrations to begin building Oculus-based content immediately.

The “Unreal Engine 3 Integration Overview” document and the “Unity Integration Overview” document, available from the Oculus Developer Center, detail the steps required to set up your UE3/Unity plus Oculus development environment.

We also recommend reading through the “Oculus Best Practices Guide”, which has tips, suggestions, and

research oriented around developing great VR experiences. Topics include control schemes, user interfaces, cut-scenes, camera features, and gameplay. The “Best Practices Guide” should be a go-to reference when designing your Oculus-ready games.

Aside from that, the next step is to get started building your own Oculus-ready games! Thousands of other developers, like you, are out there building the future of virtual reality gaming. You can reach out to them by visiting <http://developer.oculusvr.com/forums>.

5 LibOVR Integration Tutorial

If you've made it this far, you are clearly interested in integrating the Rift with your own game engine. Awesome. We are here to help.

We've designed the Oculus SDK to be as easy to integrate as possible. This section outlines a basic Oculus integration into a C++ game engine or application. We'll discuss initializing the LibOVR kernel, device enumeration, head tracking, and rendering for the Rift.

Many of the code samples below are taken directly from the OculusRoomTiny demo source code (available in `Oculus/LibOVR/Samples/OculusRoomTiny`). OculusRoomTiny and OculusWorldDemo are great places to draw sample integration code from when in doubt about a particular system or feature.

5.1 Outline of Integration Tasks

To add Oculus support to a new game, you'll need to do the following:

1. Initialize LibOVR.
2. Enumerate Oculus devices, creating HMD device and sensor objects.
3. Integrate head-tracking into your game's view and movement code. This involves:
 - (a) Reading data from the Rift's sensors through the `SensorFusion` class.
 - (b) Applying the calculated Rift orientation to the camera view, while combining it with other controls.
 - (c) Modifying movement and game play to consider head orientation.
4. Modify game rendering to integrate the HMD, including:
 - (a) Stereoscopic 3D rendering for each eye.
 - (b) Correctly computing projection, ϕ_{fov} , and other parameters based on the HMD settings.
 - (c) Applying a pixel shader to correct for optical distortion.
5. Customize UI screens to work well inside of the headset.

We'll first take a look at obtaining sensor data because it's relatively easy to set up. We'll then move on to the more involved subject of rendering.

5.2 Initialization of LibOVR

We initialize LibOVR's core by calling `System::Init`, which will configure logging and register a default memory allocator (that you can override):

```
#include "OVR.h"
using namespace OVR;
System::Init(Log::ConfigureDefaultLog(LogMask_All));
```

Note that `System::Init` must be called before any other `OVR_Kernel` objects are created, and `System::Destroy` must be called before program exit for proper cleanup. Another way to initialize the LibOVR core is to create a `System` object and let its constructor and destructor take care of initialization and cleanup, respectively. In the cases of `OculusWorldDemo` and `OculusRoomTiny`, the init and destroy calls are invoked by the `OVR_PLATFORM_APP` macro.

Once the system has been initialized, we create an instance of `OVR::DeviceManager`. This allows us to enumerate detected Oculus devices. All Oculus devices derive from the `DeviceBase` base class which provides the following functionality:

1. It supports installable message handlers, which are notified of device events.
2. Device objects are created through `DeviceHandle::CreateDevice` or more commonly through `DeviceEnumerator<>::CreateDevice`.
3. Created devices are reference counted, starting with a `RefCount` of 1.
4. A device's resources are cleaned up when it is `Released`, although its handles may survive longer if referenced.

We use `DeviceManager::Create` to create a new instance of `DeviceManager`. Once we've created the `DeviceManager`, we can use `DeviceManager::EnumerateDevices` to enumerate the detected Oculus devices. In the sample below, we create a new `DeviceManager`, enumerate available `HMDDevice` objects, and store a reference to the first active `HMDDevice` that we find.

```
Ptr<DeviceManager> pManager;
Ptr<HMDDevice> pHMD;
pManager = *DeviceManager::Create();
pHMD = *pManager->EnumerateDevices<HMDDevice>().CreateDevice();
```

We can learn more about a device by using `DeviceBase::GetDeviceInfo(DeviceInfo*info)`. The `DeviceInfo` structure is used to provide detailed information about a device and its capabilities. `DeviceBase::GetDeviceInfo` is a virtual function; therefore, subclasses such as `HMDDevice` and `SensorDevice` can provide subclasses of `DeviceInfo` with information tailored to their unique properties.

In the sample below, we read the vertical resolution, horizontal resolution, and screen size from an `HMDDevice` using `HMDDevice::GetDeviceInfo` with an `HMDInfo` object (subclass of `DeviceInfo`).

```
HMDInfo hmd;
if (pHMD->GetDeviceInfo(&hmd))
{
    MonitorName = hmd.DisplayDeviceName;
    EyeDistance = hmd.InterpupillaryDistance;
    DistortionK[0] = hmd.DistortionK[0];
    DistortionK[1] = hmd.DistortionK[1];
    DistortionK[2] = hmd.DistortionK[2];
    DistortionK[3] = hmd.DistortionK[3];
}
```

The same technique can be used to learn more about a `SensorDevice` object. Now that we have information about the `HMDDevice`, the next step is to set up rendering for the Rift.

5.3 Sensor Data and Head Tracking

The Oculus Rift hardware includes a gyroscope, accelerometer, and magnetometer. We combine the information from these sensors through a process known as *sensor fusion* to determine the orientation of the player's head in the real world, and to synchronize the player's virtual perspective in real-time.

The Rift orientation is reported as a rotation in a right-handed coordinate system, as illustrated in Figure 3.

This coordinate system uses the following axis definitions:

- Y is positive in the up direction.
- X is positive to the right.
- Z is positive heading backwards.

The rotation is maintained as a unit quaternion, but can also be reported in yaw-pitch-roll form. Positive rotation is counter-clockwise (CCW) when looking in the negative direction of each axis, and the component rotations are:

- **Pitch** is rotation around X , positive when pitching up.
- **Yaw** is rotation around Y , positive when turning left.
- **Roll** is rotation around Z , positive when tilting to the left in the XY plane.

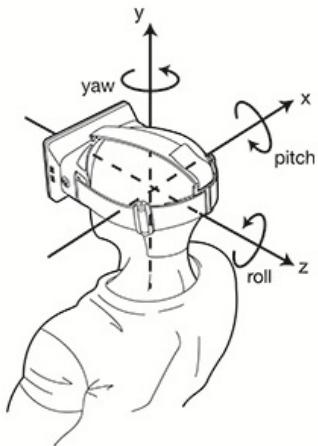


Figure 3: The Rift coordinate system

To integrate head-tracking, first we need a `SensorDevice` object to read from. If we have a reference to an HMD, we get a reference to the associated sensor using `HMDDevice::GetSensor` as follows:

```
Ptr<SensorDevice> pSensor;
pSensor = *pHMD->GetSensor();
```

We can get more information about the sensor using `SensorDevice::GetInfo`.

The `SensorFusion` class accumulates sensor notification messages to keep track of orientation. This involves integrating the gyroscope data and then using the other sensors to correct for drift. `SensorFusion` provides the orientation as a quaternion, from which users can obtain a rotation matrix or Euler angles. There are two ways to receive updates from the `SensorFusion` class:

1. We can manually pass `MessageBodyFrame` messages to the `OnMessage()` function.
2. We can attach `SensorFusion` to a `SensorDevice`. This will cause the `SensorFusion` instance to automatically handle notifications from that device.

```
SensorFusion SFusion;
if (pSensor)
    SFusion.AttachToSensor(pSensor);
```

Once an instance of `SensorFusion` is attached to a `SensorDevice`, we can use it to get relevant data from the Oculus tracker through the following functions:

```

// Obtain the current accumulated orientation.
Quatf GetOrientation() const
// Obtain the last absolute acceleration reading, in m/s^2.
Vector3f GetAcceleration() const
// Obtain the last angular velocity reading, in rad/s.
Vector3f GetAngularVelocity() const

```

In most cases, the most important data coming from SensorFusion will be the orientation quaternion Q , provided by GetOrientation. We'll make use of this to update the virtual view to reflect the orientation of the player's head. We'll also account for the orientation of the sensor in our rendering pipeline.

```

// We extract Yaw, Pitch, Roll instead of directly using the orientation
// to allow "additional" yaw manipulation with mouse/controller.
Quatf hmdOrient = SFusion.GetOrientation();
float yaw = 0.0f;
hmdOrient.GetEulerABC<Axis_Y, Axis_X, Axis_Z>(&yaw, &EyePitch, &EyeRoll);
EyeYaw += (yaw - LastSensorYaw);
LastSensorYaw = yaw;
// NOTE: We can get a matrix from orientation as follows:
Matrix4f hmdMat(hmdOrient);

```

Developers can also read the raw sensor data directly from the SensorDevice, bypassing SensorFusion entirely, by using `SensorDevice:: SetMessageHandler (MessageHandler* handler)`. The MessageHandler delegate will receive a `MessageBodyFrame` every time the tracker sends a data sample. A `MessageBodyFrame` instance provides the following data:

```

Vector3f Acceleration; // Acceleration in m/s^2.
Vector3f RotationRate; // Angular velocity in rad/s^2.
Vector3f MagneticField; // Magnetic field strength in Gauss.
float Temperature; // Temperature reading on sensor surface, in degrees Celsius.
float TimeDelta; // Time passed since last Body Frame, in seconds.

```

Over long periods of time, a discrepancy will develop between Q and the true orientation of the Rift. This problem is called *drift error*, which described more in Section 5.3.1. Errors in pitch and roll are automatically reduced by using accelerometer data to estimate the gravity vector. This feature is automatically set by default, and can be turned off by `SetGravityEnabled (false)` in the `SensorFusion` class.

Errors in yaw are reduced by magnetometer data; however, this feature is not enabled by default. For many games, such as a standard First Person Shooter (FPS), the yaw direction is frequently modified by the game controller and there is no problem. However, in many other games, the yaw error will need to be corrected. Suppose, for example, you want to maintain a cockpit directly in front of the player. It should not unintentionally drift to the side over time.

Basic use of the magnetometer is described here, with more technical details appearing in Section 5.3.2. The call `SetYawCorrectionEnabled (true)` to the `SensorFusion` enables the magnetometer to be used; however, it must first be calibrated and have a reference point stored. `OculusWorldDemo` contains examples of this usage. The object `MagCal` is initialized to maintain magnetometer calibration information. Automatic calibration is started by pressing the X key, which runs:

```

case Key_X:
    MagCal.BeginAutoCalibration(SFusion);
    SetAdjustMessage("Starting_Auto_Mag_Calibration");
    break;

```

After that, the `OnIdle` handler contains the following code to update the calibration process and detect its completion:

```
if (MagCal.IsAutoCalibrating())
{
    MagCal.UpdateAutoCalibration(SFusion);
    if (MagCal.IsCalibrated())
    {
        Vector3f mc = MagCal.GetMagCenter();
        SetAdjustMessage ("Magnetometer_Calibration_Complete\nCenter:_%f_%f_%f", mc.x, mc.y, mc.z);
    }
    SetAdjustMessage ("Mag_has_been_successfully_calibrated");
}
```

Calibration may also be performed manually, which is often more reliable. See the method `UpdateManualMagCalibration` in the `OculusWorldDemoApp` class for an example.

Once calibration has been performed, a reference point must be set, which should be a head orientation that is frequently revisited during game play. In `OculusWorldDemo`, this is set by pressing the semicolon key, resulting in the following code execution:

```
case Key_Semicolon:
    if (down)
    {
        if (SceneMode != Scene_YawView)
        {
            SFusion.SetMagReference(SFusion.GetOrientation());
            if (SFusion.IsMagReady())
                SFusion.SetYawCorrectionEnabled();
            SceneMode = Scene_YawView;
            SetAdjustMessage ("Yaw_Angle_Marks");
        }
        else
        {
            SceneMode = Scene_World;
            SetAdjustMessage ("Marks_Off");
        }
    }
    break;
```

The sample code also places marks in the field of view that illustrate how the yaw correction is working. More details appear in Section 5.3.2.

5.3.1 Some details on sensor fusion

The most important part of sensor fusion is the integration of angular velocity data from the gyroscope. In each tiny interval of time, a measurement of the angular velocity arrives:

$$\omega = (\omega_x, \omega_y, \omega_z). \quad (1)$$

Each component is the rotation rate in radians per second about its corresponding axis. For example, if the sensor were sitting on the center of a spinning hard drive disk that sits in the XZ plane, then ω_y is the rotation rate about the Y axis, and $\omega_x = \omega_z = 0$.

What happens when the Rift is rotating around some arbitrary axis? Let

$$\ell = \sqrt{\omega_x^2 + \omega_y^2 + \omega_z^2} \quad (2)$$

be the length of ω . It turns out that the angular velocity vector ω beautifully decomposes in the following useful way:

- The current axis of rotation (unit length) is the normalized gyro reading: $\frac{1}{\ell}\omega$
- The length ℓ is the rate of rotation about that axis.

Therefore, this simple piece of SDK code updates the Rift orientation Q over the interval deltaT :

```
Vector3f    rotAxis = angVel / angVelLength;
float      halfRotAngle = angVelLength * deltaT * 0.5f;
float      sinHRA   = sin(halfRotAngle);
Quatf      deltaQ(rotAxis.x*sinHRA, rotAxis.y*sinHRA, rotAxis.z*sinHRA, cos(halfRotAngle));
Q = Q * deltaQ;
```

Because of the high sampling rate of the Rift sensor and the accuracy of the gyroscope, the orientation is well maintained over a long time (several minutes). However, *drift error* eventually accumulates. Without a perfect reference orientation, the calculated Rift orientation gradually drifts away from the true orientation.

Fortunately, the other Rift sensors provide enough reference data to compensate for the drift error. The drift error can be considered as a rotation that relates the true orientation to the calculated orientation. Based on the information provided by the Rift sensors, it is helpful to decompose the drift error into two independent components:

1. **Tilt error:** An orientation change in either pitch or roll components (recall Figure 3). Imagine orientations that would cause a ball to start rolling from an otherwise level surface. A tilt error causes confusion about whether the floor is level in a game.
2. **Yaw error:** An orientation change about the Y axis only. A yaw drift error causes uncertainty about which way you are facing.

These distinctions are relative to a fixed, global coordinate system in which Y is always “up”. Because you will presumably be using the Rift on Earth, this is naturally defined as the vector outward from the center of the planet. This is sufficient for defining tilt (how far is the perceived Y from the actual Y ?). The choice of Z is arbitrary, but remains fixed for the discussion of yaw error. (X is determined automatically by the cross product of Z and Y .)

Tilt error is reduced using the accelerometer’s measurement of the gravity vector. Because the accelerometer simultaneously measures both linear head acceleration and gravity, the tilt correction method attempts to detect when head acceleration is unlikely to be occurring. In this case, it measures the difference between the measured Y axis and the calculated Y axis. When a large error is detected, tiny rotational corrections are made about a rotation axis that lies in the XZ plane and is perpendicular to both the measured and perceived Y axes. The accelerometer cannot sense yaw error because its rotation axis (Y) is aligned perfectly with gravity.

5.3.2 Magnetometer-based yaw error correction

After several minutes of game play, the Rift will gradually accumulate a few degrees of yaw error, which has been observed empirically to grow linearly over time. The magnetometer inside of the Rift has the ability to

provide yaw error correction by using the Earth’s magnetic field, which provides a fixed, 3D vector at every location. Unfortunately, this field is corrupted by interference from materials around the sensor, both inside of the Rift and in the surrounding environment. It therefore needs to be partly calibrated before it becomes useful.

Our current calibration procedure requires that magnetometer readings are taken from *four* different Rift orientations, precisely at the time and place where the user is about to play. The extracted calibration parameters are useful only for a particular Rift at a particular location. If it is moved by a meter, then the parameters may be invalid. If you want to return to the same place on another day with the same Rift, and try to use a previously stored calibration, then experiment at your own risk!

The four required magnetometer readings need to be “well separated” to provide an accurate calibration. Roughly, this means that they need to be far from each other and far from being coplanar. The SDK automatically detects whether four readings satisfy these conditions.

There are two alternatives to obtaining the four required readings before game play:

1. **Manual:** Request the player to look in a series of directions. The particular orientations are not important, but the SDK will check to make sure they are well separated.
2. **Automatic:** Simply turn on the procedure, and the calibration software will gather up samples until it finds four that are well separated.

Samples of these are provided in OculusWorldDemo by pressing the Z and X keys, respectively. The choice of which method to use is up to you as a developer. It may be burdensome to drag the user through a manual procedure, but the calibration will safely terminate. Automatic calibration is appealing because the user does not even need to be aware of the process; however, you must be certain that they will look in sufficiently distinct directions within the first couple of minutes.

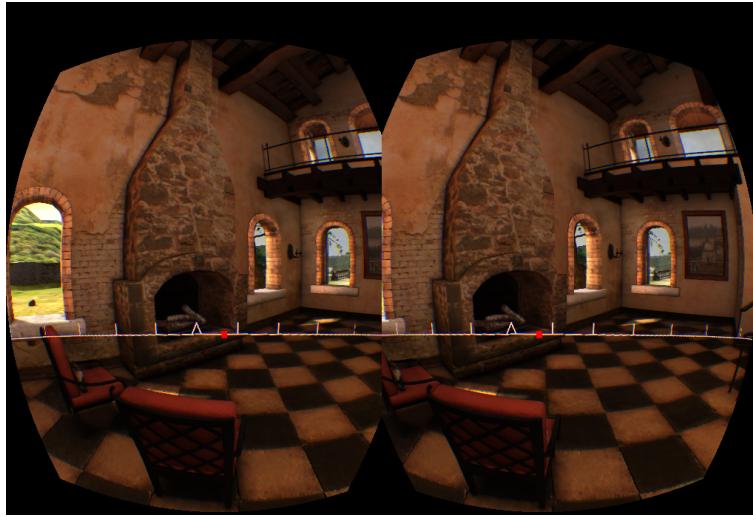


Figure 4: After pressing “;” in OculusWorldDemo, lines appear showing the direction of the reference point. When the Rift orientation is close to the reference point (triangle in the center), then the sensor fusion may detect yaw error (assuming calibration was first performed). A small square shows the forward looking direction. If the square is green, then yaw correction is taking place (and you might notice a slow yaw rotation); otherwise, it is red.

Once the calibration is completed, a *reference point* needs to be selected. This should be an orientation that is frequently revisited, such as looking forward and level. In `OculusWorldDemo`, press “;” to set the reference point and visualize its location. Once the magnetometer has been calibrated and the reference point has been set, the sensor fusion module will attempt to detect and correct yaw errors whenever the Rift orientation is somewhat close to the reference point. Figure 4 shows how to determine whether the yaw correction is working properly in `OculusWorldDemo`.

To use the magnetometer for yaw correction, you should create a `MagCalibration` object, which is defined in `Util_Mag_Calibration.h`. The methods allow a manual or automatic calibration procedure to be quickly integrated into your software. See example usage in `OculusWorldDemo.cpp`, with references to the `MagCal` object, which is a protected member of the `OculusWorldDemoApp` class. The class also contains `UpdateManualMagCalibration()`, which is a customized method for performing manual calibration.

If you are adventurous, you might want to customize the magnetometer calibration and yaw correction. For example, you could set smaller thresholds for declaring the points to be well-separated. This would be accomplished in `OculusWorldDemo` as:

```
CalMag.SetMinMagDistance(0.2f);  
CalMag.SetMinQuatDistance(0.4f);
```

This overrides their default values of 0.3 and 0.5. Their units are Gauss and Euclidean distance in normalized quaternion coordinates, respectively. The lower values could make the automatic calibration procedure achieve its goal more easily, but comes at the risk of poorer estimation of the calibration parameters. You can also change parameters of the `SensorFusion` object. For example, using `SetMagRefDistance(0.25f)` will enlarge the region over which yaw drift correction will be activated around the reference point, but it could make an incorrect assessment about the yaw error. A significant improvement would be to store a sparse cloud of reference points, to ensure that near enough references are always available. This and related enhancements could come in future releases.

5.4 User Input Integration

Head tracking will need to be integrated with an existing control scheme for many games to provide the most comfortable, intuitive, and usable interface for the player.

For example, in an FPS, the player moves forward, backward, left, and right using the left joystick, and looks left, right, up, and down using the right joystick. When using the Rift, the player can now look left, right, up, and down, using their head. However, players should not be required to frequently turn their heads 180 degrees; they need a way to reorient themselves so that they are always comfortable (the same way we turn our bodies if we want to look behind ourselves for more than a brief glance).

As a result, developers should carefully consider their control schemes and how to integrate head-tracking when designing games for VR. The `OculusRoomTiny` application provides a source code sample for integrating Oculus head tracking with the aforementioned, standard FPS control scheme.



Figure 5: OculusWorldDemo stereo rendering.

5.5 Rendering Configuration

As you may be aware, Oculus rendering requires split-screen stereo with distortion correction for each eye to account for the Rift’s optics. Setting this up can be tricky, but immersive rendering is what makes the Rift magic come to life. We separate our rendering description into several sections:

- Section 5.5.1 introduces the basics of HMD stereo rendering and projection setup.
- Section 5.5.2 covers distortion correction, describing the pixel shader and its associated parameters.
- Sections 5.5.3 and 5.5.4 round out the discussion by explaining the scaling and field of view math necessary for the scene to look correct.
- Finally, Section 5.5.5 introduces the `StereoConfig` utility class that does a lot of the hard work for you, hiding math complexity behind the scenes.

Aside from changes to the game engine’s rendering pipeline, 60 FPS low latency rendering is also critical for immersive VR. We cover VSync, latency, and other performance requirements in Section 5.5.6.

5.5.1 Rendering stereo

The Oculus Rift requires the game scene to be rendered in split-screen stereo, with half the screen used for each eye. When using the Rift, your left eye sees the left half of the screen, whereas the right eye sees the right half. This means that your game will need to render the entire scene twice, which can be achieved with logic similar to the following pseudo code:

```
// Render Left Eye Half
SetViewport(0, 0, HResolution/2, VResolution);
SetProjection(LeftEyeProjectionMatrix);
RenderScene();

// Render Right Eye Half
SetViewport(HResolution/2, 0, HResolution, VResolution);
SetProjection(RightEyeProjectionMatrix);
RenderScene();
```

Note that the *reprojection stereo rendering* technique, which relies on left and right views being generated from a single fully rendered view, is not usable inside of an HMD because of significant artifacts at object edges.

Unlike stereo TVs, rendering inside of the Rift does not require off-axis or asymmetric projection. Instead, projection axes are parallel to each other as illustrated in Figure 6. This means that camera setup will be very similar to that normally used for non-stereo rendering, except you will need to shift the camera to adjust for each eye location.

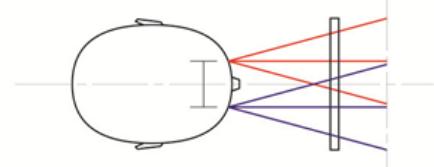


Figure 6: HMD eye view cones.

To get correct rendering on the Rift, the game needs to use the physically appropriate field of view ϕ_{fov} , calculated based on the Rift's dimensions. The parameters needed for stereo rendering are reported from LibOVR in OVR::HMDInfo as follows:

Member Name	Description
HScreenSize, VScreenSize	Physical dimensions of the entire HMD screen in meters. Half HScreenSize is used for each eye. The current physical screen size is 149.76 x 93.6mm, which will be reported as 0.14976f x 0.0935f.
VScreenCenter	Physical offset from the top of the screen to eye center, in meters. Currently half VScreenSize.
EyeToScreenDistance	Distance from the eye to the screen, in meters. This combines distances from the eye to the lens, and from the lens to the screen. This value is needed to compute the ϕ_{fov} correctly.
LensSeparationDistance	Physical distance between the lens centers, in meters. Lens centers are the centers of distortion; we will talk about them later in Section 5.5.2
InterpupillaryDistance	Configured distance between eye centers.
HResolution, VResolution	Resolution of the entire HMD screen in pixels. Half the HResolution is used for each eye. The reported values are 1280 × 800 for the DK, but we are determined to increase this in the future!
DistortionK	Radial distortion correction coefficients, discussed in Section 5.5.2.

So, how do we use these values to set up projection? For simplicity, let us focus on rendering for the left eye and ignore the distortion for the time being. Before you can draw the scene you'll need to take several steps:

1. Set the viewport to cover the left eye screen area.
2. Determine the aspect ratio a and ϕ_{fov} based on the reported HMDInfo values.
3. Calculate the center projection matrix \mathbf{P} based on a and ϕ_{fov} .
4. Adjust the projection matrix \mathbf{P} based on lens separation distance.
5. Adjust the view matrix \mathbf{V} to match eye location.

Setting up the viewport is easy, simply set it to `(0, 0, HResolution/2, VResolution)` for the left eye. In most 3D graphics systems, the clip coordinates [-1,1] will be mapped to fill the viewport, with (0,0)

corresponding to the center of projection.

Ignoring distortion, Rift half-screen aspect ratio a and vertical FOV ϕ_{fov} are determined by

$$a = \frac{\text{HResolution}}{2 \cdot \text{VResolution}} \quad (3)$$

and

$$\phi_{fov} = 2 \arctan \left(\frac{\text{VScreenSize}}{2 \cdot \text{EyeToScreenDistance}} \right). \quad (4)$$

We form the projection matrix, \mathbf{P} , based on a and ϕ_{fov} , as

$$\mathbf{P} = \begin{bmatrix} \frac{1}{a \cdot \tan(\phi_{fov}/2)} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\phi_{fov}/2)} & 0 & 0 \\ 0 & 0 & \frac{z_{far}}{z_{near} - z_{far}} & \frac{z_{far} z_{near}}{z_{near} - z_{far}} \\ 0 & 0 & -1 & 0 \end{bmatrix}, \quad (5)$$

in which z_{near} and z_{far} are the standard clipping plane depth coordinates. This common calculation can be done by the `Matrix4f::PerspectiveRH` function in the Oculus SDK, the `gluPerspective` utility function in OpenGL, or `D3DXMatrixPerspectiveFovRH` in Direct3D.

The projection center of \mathbf{P} as computed above falls in the center of each screen, so we need to modify it to coincide with the center of the lens instead. Here, lens separation is used instead of the eye separation due to the properties of the collimated light. Lens center adjustment can be done in final clip coordinates, computing the final left and right projection matrices as illustrated below. Let h denote the absolute value of horizontal offset to account for lens separation. This can be used in a transformation matrix

$$\mathbf{H} = \begin{bmatrix} 1 & 0 & 0 & \pm h \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \quad (6)$$

which is applied at the end to obtain $\mathbf{P}' = \mathbf{H} \mathbf{P}$. In the upper right corner of \mathbf{H} , the term h appears for the left-eye case, and $-h$ appears for the right eye. For convenience, let

$$d_{lens} = \text{LensSeparationDistance}. \quad (7)$$

The particular horizontal shift in meters is

$$h_{meters} = \frac{\text{HScreenSize}}{4} - \frac{d_{lens}}{2}. \quad (8)$$

In screen coordinates,

$$h = \frac{4 h_{meters}}{\text{HScreenSize}}. \quad (9)$$

In terms of screen size, this adjustment is significant: Assuming 64mm interpupillary distance (IPD) and 149.76mm screen size of the 7" Rift, each eye projection center needs to be translated by about 5.44mm towards the center of the device. This is a critical step for correct stereo rendering.

Assuming that the original non-stereo game view transform \mathbf{V} falls at the center between the eyes, the final adjustment we have to make is shift that view horizontally to match each eye location:

$$\mathbf{V}' = \begin{bmatrix} 1 & 0 & 0 & \pm d_{ip}/2 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \mathbf{V}. \quad (10)$$

where

$$d_{ip} = \text{InterpupillaryDistance}. \quad (11)$$

It is important that this shift is done by half of the interpupillary distance in world units. Please refer to your game's content or design documents for the conversion to or from real-world units. In Unreal Engine 3, for example, 2 units = 1 inch. In Unity, 1 unit = 1 meter.

We can now present a more complete example of this stereo setup, as it is implemented inside of the OVR::Util::Render::StereoConfig utility class. It covers all of the steps described in this section with exception of viewport setup.

```
HMDInfo& hmd = ...;
Matrix4f viewCenter = ...;

// Compute Aspect Ratio. Stereo mode cuts width in half.
float aspectRatio = float(hmd.HResolution * 0.5f) / float(hmd.VResolution);

// Compute Vertical FOV based on distance.
float halfScreenDistance = (hmd.VScreenSize / 2);
float yfov = 2.0f * atan(halfScreenDistance/HMD.EyeToScreenDistance);

// Post-projection viewport coordinates range from (-1.0, 1.0), with the
// center of the left viewport falling at (1/4) of horizontal screen size.
// We need to shift this projection center to match with the lens center.
// We compute this shift in physical units (meters) to correct
// for different screen sizes and then rescale to viewport coordinates.
float viewCenter           = hmd.HScreenSize * 0.25f;
float eyeProjectionShift   = viewCenter - hmd.LensSeparationDistance*0.5f;
float projectionCenterOffset = 4.0f * eyeProjectionShift / hmd.HScreenSize;

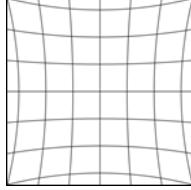
// Projection matrix for the "center eye", which the left/right matrices are based on.
Matrix4f projCenter = Matrix4f::PerspectiveRH(yfov, aspect, 0.3f, 1000.0f);
Matrix4f projLeft   = Matrix4f::Translation(projectionCenterOffset, 0, 0) * projCenter;
Matrix4f projRight  = Matrix4f::Translation(-projectionCenterOffset, 0, 0) * projCenter;

// View transformation translation in world units.
float halfIPD = hmd.InterpupillaryDistance * 0.5f;
Matrix4f viewLeft = Matrix4f::Translation(halfIPD, 0, 0) * viewCenter;
Matrix4f viewRight= Matrix4f::Translation(-halfIPD, 0, 0) * viewCenter;
```

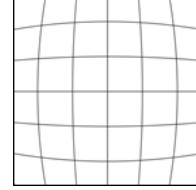
With all of this setup done, you should be able to see the 3D world and converge on it inside of the Rift. The last, and perhaps more challenging step, will be correcting for distortion due to the lenses.

5.5.2 Distortion correction

The lenses in the Rift magnifies the image to provide an increased FOV, but this comes at the expense of creating a pincusion distortion of the image.



Pincushion Distortion



Barrel Distortion

This radially symmetric distortion can be corrected in software by introducing a barrel distortion that cancels it out. Our distortion models are expressed in terms of the distance r from the center (the axis of symmetry). It is helpful to think in polar coordinates (r, θ) for a point in the image. A distortion model performs the following transformation:

$$(r, \theta) \mapsto (f(r) r, \theta), \quad (12)$$

in which the *scaling function* is given by the following standard model:

$$f(r) = k_0 + k_1 r^2 + k_2 r^4 + k_3 r^6. \quad (13)$$

In other words, the direction θ is unharmed, but the distance r is expanded or contracted. The four coefficients k_0, \dots, k_3 control the distortion. They are all positive in the case of barrel distortion. With properly chosen coefficients, the pincushion distortion is canceled off by a barrel distortion.

In OculusWorldDemo, this is implemented by the Direct3D10 pixel shader shown in Figure 7.

This shader is designed to run on a rectangle covering half of the screen, while the input texture spans both the left and right eyes. The input texture coordinates, passed in as `oTexCoord`, range from (0,0) for the top left corner of the Oculus screen, to (1,1) at the bottom right. This means that for the left eye viewport, `oTexCoord` will range from (0,0) to (0.5,1). For the right eye it will range from (0.5,0) to (1,1).

The distortion function used by `HmdWarp` is, however, designed to operate on [-1,1] unit coordinate range, from which it can compute the radius. This means that there are a number of variables needed to scale and center the coordinates properly to apply the distortion. These are:

Variable Name	Description
<code>ScaleIn</code>	Rescale input texture coordinates to [-1,1] unit range, and correct aspectratio.
<code>Scale</code>	Rescale output (sample) coordinates back to texture range and increase scale so as to support sampling outside of the screen.
<code>HmdWarpParam</code>	The array of distortion coefficients (<code>DistortionK[]</code>).
<code>ScreenCenter</code>	Texture coordinate for the center of the half-screen texture. This is used to clamp sampling, preventing pixel leakage from one eye view to the other.
<code>LensCenter</code>	Shifts texture coordinates to center the distortion function around the center of the lens.

```

Texture2D    Texture : register(t0);
SamplerState Linear : register(s0);
float2       LensCenter;
float2       ScreenCenter;
float2       Scale;
float2       ScaleIn;
float4       HmdWarpParam;

// Scales input texture coordinates for distortion.
float2 HmdWarp(float2 in01)
{
    float2 theta = (in01 - LensCenter) * ScaleIn; // Scales to [-1, 1]
    float rSq = theta.x * theta.x + theta.y * theta.y;
    float2 rvector= theta * (HmdWarpParam.x + HmdWarpParam.y * rSq +
                            HmdWarpParam.z * rSq * rSq +
                            HmdWarpParam.w * rSq * rSq * rSq);
    return LensCenter + Scale * rvector;
}
float4 main(in float4 oPosition : SV_Position, in float4 oColor : COLOR,
            in float2 oTexCoord : TEXCOORD0) : SV_Target
{
    float2 tc = HmdWarp(oTexCoord);
    if (any(clamp(tc, ScreenCenter-float2(0.25,0.5),
                  ScreenCenter+float2(0.25, 0.5)) - tc))
        return 0;
    return Texture.Sample(Linear, tc);
};

```

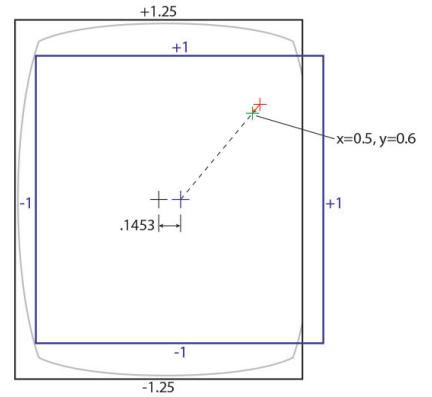
Figure 7: Pixel shader code for removing lens distortion.

The final variable is calculated as

$$\begin{aligned}
 \text{LensCenter} &= 4 \left(\frac{\frac{\text{HScreenSize}}{4} - \frac{\text{LensSeparationDistance}}{2}}{\frac{\text{HScreenSize}}{2}} \right) \\
 &= 1 - 2 \cdot \text{LensSeparationDistance} / \text{HScreenSize}.
 \end{aligned} \tag{14}$$

The following diagram illustrates the left eye distortion function coordinate range, shown as a blue rectangle, as it relates to the left eye viewport coordinates. As you can see, the center of distortion has been shifted to the right in relation to the screen center to align it with axis through the center of the lense. For the 7" screen and 64mm lense separation distance, viewport shift is roughly 0.1453 coordinate units. These parameters may change for future headsets and so this should always be computed dynamically.

The diagram also illustrates how sampling coordinates are mapped by the distortion function. A distortion unit coordinate of $(0.5, 0.6)$ is marked as a green cross; it has a radius of ≈ 0.61 . In the example shown, this maps to a sampling radius of ≈ 0.68 post-distortion, illustrated by a red cross. As a result of the distortion shader, pixels in the rendered image move towards the center of distortion, or from red to green in the diagram. The amount of displacement increases the further out we go from the distortion center.



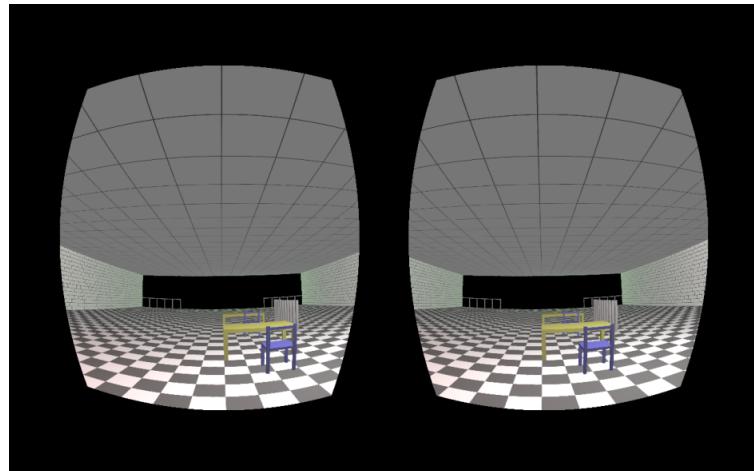
Hopefully, it's clear from this discussion that the barrel distortion pixel shader needs to run as a post-process on the rendered game scene image. This has several implications:

- The original scene rendering will need to be done to a render target.
- The scene render target will need to be larger than the final viewport, to account for the distortion pulling pixels in towards the center.
- The FOV and image scale will need to be adjusted do accommodate for the distortion.

We will now discuss the distortion scale, render target, and FOV adjustments necessary to make the image look correct inside of the Rift.

5.5.3 Distortion scale

If you run the distortion shader on the original image render target that is the same size as the output screen you will get an image similar to the following:



Here, the pixels at the edges have been pulled in towards the center, with black being displayed outside, where no pixel data was available. Although this type of rendering would look acceptable within the Rift,

a significant part of the FOV is lost because large areas of the screen go unused. How would we make the scene fill up the entire screen?

The simplest solution is to increase the scale of the input texture, controlled by the `Scale` variable of the distortion pixel shader discussed earlier. As an example, if we want to increase the perceived input texture size by 25% we can adjust the sampling coordinate `Scale` by a factor of $(1/1.25) = 0.8$. Doing so will have several effects:

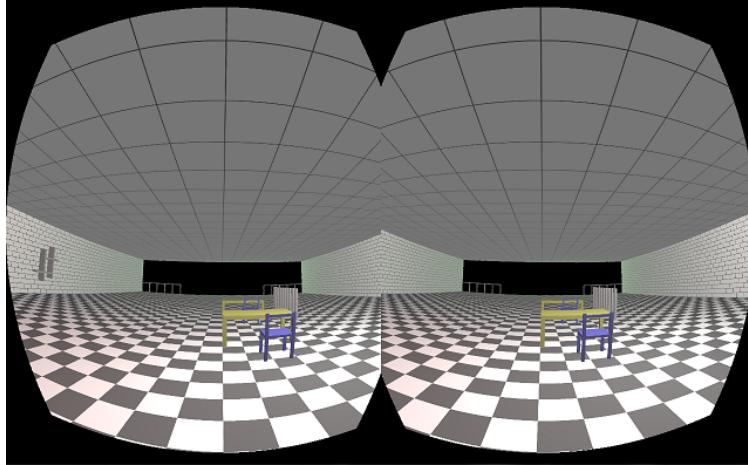
- The size of the post-distortion image will increase on screen.
- The required rendering FOV will increase.
- The quality of the image will degrade due to sub-sampling from the scaled image, resulting in blocky or blurry pixels around the center of the view.

Since we really don't want the quality to degrade, the size of the source render target can be increased by the same amount to compensate. For the 1280×800 resolution of the Rift, a 25% scale increase will require rendering a 1600×1000 buffer. Unfortunately, this incurs a 1.56 times increase in the number of pixels in the source render target. However, we don't need to completely fill the far corners of the screen where the user cannot see. Trade-offs are evident between the covered field of view, quality, and rendering performance.

For the 7" Rift, we recommend selecting the radius (distance from image center) so that the image reaches the left screen edge. For example, to fit the left side of the display $(-1, 0)$:

$$s = f(-1 - \text{LensCenter}), \quad (15)$$

in which s is the scaling factor. This yields maximal horizontal FOV without filling the pixels at the top of the screen, which are not visible to most users:



For `OculusWorldDemo`, the actual distortion scale factor is computed inside of the `StereoConfig::updateDistortionOffsetAndScale` function by fitting the distortion radius to the left edge of the viewport. The `LensCenter` is the same as used in the shader above.

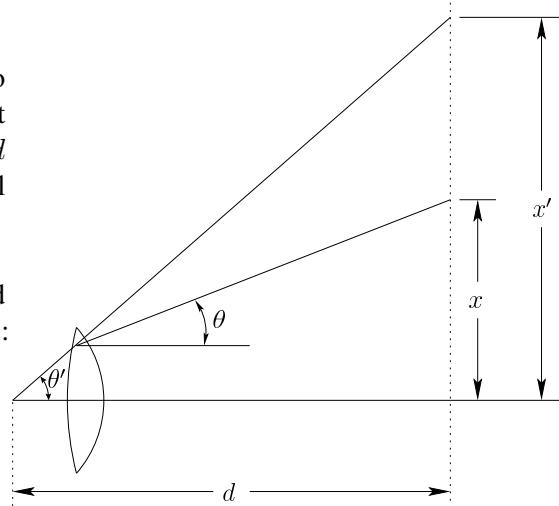
5.5.4 Distortion and FOV

With distortion scale and offset properly computed, the FOV needs to be corrected. For this, we need to examine the geometry of the Rift projection as illustrated in the diagram on the right.

The diagram illustrates the effect that a lens introduces into an otherwise simple calculation of the FOV, assuming that the user is looking at the screen through the lens. Here, d specifies the *eye to screen* distance and x is half the vertical screen size ($x = \text{vScreenSize}/2$).

In the absence of the lens, the FOV can be easily computed based on the distances d and x , as described in Section 5.5.1:

$$\phi_{fov} = 2 \arctan(x/d)$$



The lens, however, increases the perceived screen height from $2x$ to $2x'$, in which x' can be computed through the scaling function (13). Thus, the effective FOV is obtained as

$$\phi'_{fov} = 2 \arctan(x'/d) \quad (16)$$

in which $x' = s \cdot \text{vScreenSize}/2$ and s is the scaling factor from Section 5.5.3 (recall Equation 15).

Note that we compute the field of view of the distorted render target (RT), which is affected by the distortion scale s and may not necessarily match the screen edge. Assuming that both the RT and the display have the same aspect ratios, it is sufficient to adjust the screen size to correctly obtain ϕ'_{fov} .

5.5.5 StereoConfig utility class

If setting up the projection and distortion scaling seems like a lot of work, you'll be happy to learn that the Oculus SDK comes with a set of full-source utility classes that do a lot of the work for you. The most important class for rendering setup is `StereoConfig`, located inside of the `OVR::Util::Render` namespace. This class works as follows:

1. First, create an instance of `StereoConfig` and initialize it with `HMDInfo`, viewport, distortion fit location, IPD and other desired settings.
2. `StereoConfig` computes the rendering scale, distortion offsets, FOV and projection matrices.
3. `StereoConfig::GetEyeRenderParams` returns the projection matrix and distortion settings for each eye. These can be used for rendering stereo inside of the game loop.

`StereoConfig` class is used for rendering setup inside of the `OculusRoomTiny` and `OculusWorldDemo` samples. Here's an example of the initialization you need to get started:

```
using namespace OVR::Util::Render;

HMDDevice* pHMD = ...;
StereoConfig stereo;
HMDInfo hmd;
float renderScale;

// Obtain setup data from the HMD and initialize StereoConfig
```

```

// for stereo rendering.
pHMD->GetDeviceInfo(&hmd);

stereo.SetFullViewport (Viewport (0,0, Width, Height));
stereo.SetStereoMode (Stereo_LeftRight_Multipass);
stereo.SetHMDInfo (hmd);
stereo.SetDistortionFitPointVP (-1.0f, 0.0f);

renderScale = stereo.GetDistortionScale();

```

As you can see, after all parameters are initialized, `GetDistortionScale` computes the rendering scale that should be applied to the render texture. This is the scale that will maintain one-to-one rendering quality at the center of the screen while simultaneously scaling the distortion to fit its left edge.

Based on this computed state, you can get left and right eye rendering parameters as follows:

```

StereoEyeParams leftEye = stereo.GetEyeRenderParams (StereoEye_Left);
StereoEyeParams rightEye = stereo.GetEyeRenderParams (StereoEye_Right);

// Left eye rendering parameters
Viewport      leftVP       = leftEye.VP;
Matrix4f      leftProjection = leftEye.Projection;
Matrix4f      leftViewAdjust = leftEye.ViewAdjust;

```

You can use the resulting `Viewport` and projection matrix directly for rendering the scene. `ViewAdjust` should be a post-transform applied after the game's view matrix to properly shift the camera for the left or right eye.

5.5.6 Rendering performance

Aside from changes to the game engine's renderer to account for the Rift's optics, there are two other requirements when rendering for VR:

- The game engine should run at least 60 frames per second without dropping frames.
- Vertical Sync (vsync) should always be enabled to prevent the player from seeing screen tearing.

These may seem arbitrary, but our experiments have shown them to be important for a good VR experience. A player can easily tell the difference between 30 FPS and 60 FPS when playing in VR because of the immersive nature of the game. The brain can suspend disbelief at 60 FPS. At 30 FPS, the world feels choppy. Vertical sync is also critical. Since the Rift screen covers all of the player's view, screen tearing is very apparent, and causes artifacts that break immersion.

6 Optimization

6.1 Latency

Minimizing latency is crucial to immersive VR and low latency head tracking is part of what sets the Rift apart. We define latency as the time between movement of the player's head, and the updated image being displayed on the screen. We call this latency loop "motion-to-photon" latency. The more you can minimize motion-to-photon latency in your game, the more immersive the experience will be for the player.

Two other important concepts are actual latency and perceived latency.

Actual latency is equivalent to motion-to-photon latency. It is the latency in the system at the hardware and software level.

Perceived latency is how much latency the player perceives when using the headset. Perceived latency may be less than actual latency depending on the player's movements and by employing certain techniques in software.

We're always working to reduce actual and perceived latency in our hardware and software pipeline. For example, in some cases we're able to reduce perceived latency by 20ms or more using a software technique called predictive tracking.

Although 60ms is a widely cited threshold for acceptable VR, at Oculus we believe the threshold for compelling VR to be below 40ms of latency. Above this value you tend to feel significantly less immersed in the environment. Obviously, in an ideal world, the closer we are to 0ms, the better.

For the Rift developer kit, we expect the actual latency to be approximately 30ms to 50ms. This depends partly on the screen content. For example, a change from black to dark brown may take 5ms but a larger change in color from black to white may take 20ms.

Stage	Event	Event Duration	Worst Case Total
Start	Oculus tracker sends data	N/A	0ms
Transit	Computer receives tracker data	≈ 2ms	≈ 2ms
Processing	Game engine renders latest frame (60 FPS w/ vysnc)	≈ 0 to 16.67ms	≈ 19ms
Processing	Display controller writes latest frame to LCD (top to bottom)	≈ 16.67ms	≈ 36ms
Processing	Simultaneously, pixels switching colors	≈ 0 to 15ms	≈ 51ms
End	Latest frame complete; presented to user	N/A	≈ 51ms

Again, these numbers represent the actual latency assuming a game running at 60 FPS with vsync enabled. Actual latency will vary depending on the scene being rendered. Perceived latency can be reduced further. As developers, we want to do everything we can to reduce latency in this pipeline. Techniques for Reducing Latency:

- Run at 60 FPS (remember that vsync should always be enabled for VR).
- Minimize swap-chain buffers to a maximum of 2 (the on screen and off screen buffers).
- Reduce the amount of rendering work where possible. Multi-pass rendering and complex shaders

increase the rendering latency and hence the time between reading the HMD orientation and having the frame ready to display.

- Reduce render command buffer size. By default the driver may buffer several frames of render commands to batch GPU transfers and smooth out variability in rendering times. This needs to be minimized. One technique is to make a rendering call that blocks until the current frame is complete. This can be a “block until render queue empty event” or a command that reads back a property of the rendered frame. While blocking, we’re preventing additional frames from being submitted and hence buffered in the command queue.

A Display Device Management

A.1 Display Identification

Display devices identify themselves and their capabilities using EDID¹. When the device is plugged into a PC, the display adapter reads a small packet of data from it. This includes the manufacturer code, device name, supported display resolutions, and information about video signal timing. When running an OS that supports multiple monitors, the display is identified and added to a list of active display devices which can be used to show the desktop or fullscreen applications.

The display within the Oculus Rift interacts with the system in the same way as a typical PC monitor. It too provides EDID information which identifies it as having a manufacturer code of ‘OVR’, a model ID of ‘Rift DK1’, and support for several display resolutions including its native 1280×800 at 60Hz.

A.2 Display Configuration

After connecting a Rift to the PC it is possible to modify the display settings through the Windows Control Panel², or the System Preferences Display panel on MacOS.

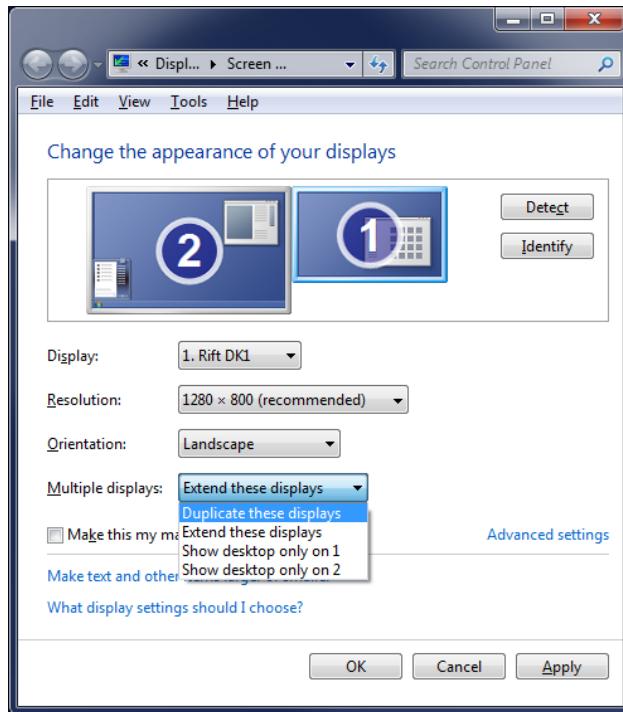


Figure 8: Screenshot of the Windows “Screen Resolution” dialog.

Figure 8 shows the “Screen Resolution” dialog for a PC with the Rift display and a PC monitor connected. In this configuration there are four modes that can be selected as shown in the figure. These are duplicate mode, extended mode, and standalone mode for either of the displays.

¹Extended Display Identification Data

²Under Windows 7 this can be accessed via Control Panel : All Control Panel Items : Display : Screen Resolution

A.2.1 Duplicate display mode

In duplicate display mode the same portion of the desktop is shown on both displays, and they adopt the same resolution and orientation settings. The OS attempts to choose a resolution which is supported by both displays, while favoring the native resolutions described in the EDID information reported by the displays. Duplicate mode is a potentially viable mode in which to configure the Rift, however it suffers from vsync issues. See Section A.4 for details.

A.2.2 Extended display mode

In extended mode the displays show different portions of the desktop. The Control Panel can be used to select the desired resolution and orientation independently for each display. Extended mode suffers from shortcomings related to the fact that the Rift is not a viable way to interact with the desktop, nevertheless it is the current recommended configuration option. The shortcomings are discussed in more detail in Section A.4.

A.2.3 Standalone display mode

In standalone mode the desktop is displayed on just one of the plugged in displays. It is possible to configure the Rift as the sole display, however it becomes impractical due to problems interacting with the desktop.

A.3 Selecting A Display Device

Reading of EDID information from display devices has been found to occasionally be slow and unreliable. In addition, EDID information may be cached leading to problems with stale data. As a result, display devices may sometimes become associated with incorrect display names and resolutions, with arbitrary delays before the information becomes current.

Because of these issues we adopt an approach which attempts to identify the Rift display name among the attached display devices, however we do not require that it be found for an HMD device to be created using the API.

If the Rift display device is not detected but the Rift is detected through USB, then we return an empty display name string. In this case your application could attempt to locate it using additional information, for example display resolution. The GetDeviceInfo call in the API returns the native resolution of the Rift display. This can be used to match one of the available display devices.

In general, due to the uncertainty associated with identifying the Rift display device, it may make sense to incorporate functionality into your application to allow the user to choose the display manually, for example from a drop-down list of enumerated display devices. One additional root cause of the above scenario, aside from EDID issues, is that the user failed to plug in the Rift video cable. We recommend appropriate assistance inside your application to help users troubleshoot an incorrectly connected Rift device.

A.3.1 Windows

The following code outlines how to create an HMD device and get the display device name on Windows:

```

DeviceManager* pManager = DeviceManager::Create();

HMDDevice* pHMD = pManager->EnumerateDevices<HMDDevice>().CreateDevice();

HMDInfo info;
pHMD->GetDeviceInfo(&HMDInfo);

char* displayName = info.DisplayDeviceName;

```

If there are no EDID issues and we detect the Rift display device successfully, then we return the display name corresponding to that device, for example \\.\DISPLAY1\Monitor0.

To display the video output on the Rift, the application needs to match the display name determined above to an object used by the rendering API. Unfortunately, each rendering system used on Windows (OpenGL, Direct3D 9, Direct3D 10-11) uses a different approach. OpenGL uses the display device name returned in HMDInfo, while in Direct3D the display name must be matched against a list of displays returned by D3D. The DesktopX and DesktopY members of HMDInfo can be used to position a window on the Rift if you do not want to use fullscreen mode.

When using Direct3D 10 or 11, the following code shows how to obtain an IDXGIOOutput interface using the DXGI API:

```

IDXGIOOutput* searchForOculusDisplay(char* oculusDisplayName)
{
    IDXGIFactory* pFactory;
    CreateDXGIFactory(__uuidof(IDXGIFactory), (void**)(&pFactory));

    UInt32 adapterIndex = 0;
    IDXGIAdapter* pAdapter;

    // Iterate through adapters.
    while (pFactory->EnumAdapters(adapterIndex, &pAdapter) != DXGI_ERROR_NOT_FOUND)
    {
        UInt32 outputIndex = 0;
        IDXGIOOutput* pOutput;

        // Iterate through outputs.
        while (pAdapter->EnumOutputs(outputIndex, &pOutput) != DXGI_ERROR_NOT_FOUND)
        {
            DXGI_OUTPUT_DESC outDesc;
            pOutput->GetDesc(&outDesc);
            char* outputName = outDesc.DeviceName;

            // Try and match the first part of the display name.
            // For example an outputName of "\\.\DISPLAY1" might
            // correspond to a displayName of "\\.\DISPLAY1\Monitor0".
            // If two monitors are setup in 'duplicate' mode then they will
            // have the same 'display' part in their display name.
            if (strstr(oculusDisplayName, outputName) == oculusDisplayName)
            {
                return pOutput;
            }
        }
    }

    return NULL;
}

```

After you've successfully obtained the IDXGIOOutput interface, you can set your Direct3D swap-chain to render to it in fullscreen mode using the following:

```
IDXGIOOutput* pOculusOutputDevice = searchForOculusDisplay(oculusDisplayName);
pSwapChain->SetFullscreenState(TRUE, pOculusOutputDevice);
```

When using Direct3D 9, you must create the Direct3DDevice with the “adapter” number corresponding to the Rift (or other display that you want to output to). This next function shows how to find the adapter number:

```
unsigned searchForOculusDisplay(const char* oculusDisplayName)
{
    for (unsigned Adapter=0; Adapter < Direct3D->GetAdapterCount(); Adapter++)
    {
        MONITORINFOEX Monitor;
        Monitor.cbSize = sizeof(Monitor);
        if (::GetMonitorInfo(Direct3D->GetAdapterMonitor(Adapter), &Monitor) && Monitor.szDevice[0])
        {
            DISPLAY_DEVICE DispDev;
            memset(&DispDev, 0, sizeof(DispDev));
            DispDev.cb = sizeof(DispDev);
            if (::EnumDisplayDevices(Monitor.szDevice, 0, &DispDev, 0))
            {
                if (strstr(DispDev.DeviceName, oculusDisplayName))
                {
                    return Adapter;
                }
            }
        }
    }
    return D3DADAPTER_DEFAULT;
}
```

Unfortunately, you must re-create the device to switch fullscreen displays. You can use a fullscreen device for windowed mode without recreating it (although it still has to be reset).

A.3.2 MacOS

When running on MacOS, the HMDInfo contains a CGDirectDisplayId which you should use to target the Rift in fullscreen. If the Rift is connected (and detected by libovr), its DisplayId is stored, otherwise DisplayId contains 0. The following is an example of how to use this to make a Cocoa NSView fullscreen on the Rift:

```
DeviceManager* pManager = DeviceManager::Create();

HMDDevice* pHMD = pManager->EnumerateDevices<HMDDevice>().CreateDevice();

HMDInfo hmdInfo;
pHMD->GetDeviceInfo(&hmdInfo);

CGDirectDisplayId RiftDisplayId = (CGDirectDisplayId) hmdInfo.DisplayId;

NSScreen* usescreen = [NSScreen mainScreen];
NSArray* screens = [NSScreen screens];
for (int i = 0; i < [screens count]; i++)
{
    NSScreen* s = (NSScreen*) [screens objectAtIndex:i];
    CGDirectDisplayID disp = [NSView displayFromScreen:s];
```

```

    if (disp == RiftDisplayId)
        usescreen = s;
}

[View enterFullScreenMode:usescreen withOptions:nil];

```

The DesktopX and DesktopY members of HMDInfo can be used to position a window on the Rift if you do not want to use fullscreen mode.

A.4 Rift Display Considerations

There are several considerations when it comes to managing the Rift display on a desktop OS.

A.4.1 Duplicate mode VSync

In duplicate monitor mode it is common for the supported video timing information to be different across the participating monitors, even when displaying the same resolution. When this occurs the video scans on each display will be out of sync and the software vertical sync mechanism will be associated with only one of the displays. In other words, swap-chain buffer switches (for example following a glSwapBuffers or Direct3D Present call) will only occur at the correct time for one of the displays, and ‘tearing’ will occur on the other display. In the case of the Rift, tearing is very distracting, and so ideally we’d like to force it to have vertical sync priority. Unfortunately, the ability to do this is not something that we’ve found to be currently exposed in system APIs.

A.4.2 Extended mode problems

When extended display mode is used in conjunction with the Rift, the desktop will extend partly onto the regular monitors and partly onto the Rift. Since the Rift displays different portions of the screen to the left and right eyes, it is not suited to displaying the desktop in a usable form, and confusion may arise if icons or windows find their way onto the portion of the desktop displayed on the Rift.

A.4.3 Observing Rift output on a monitor

Sometimes it’s desirable to be able to see the same video output on the Rift and on an external monitor. This can be particularly useful when demonstrating the device to a new user, or during application development. One way to achieve this is through the use of duplicate monitor mode as described above, however we don’t currently recommend this approach due to the vertical sync priority issue. An alternative approach is through the use of a DVI or HDMI splitter. These take the video signal coming from the display adapter and duplicate it such that it can be fed to two or more display devices simultaneously. Unfortunately this can also cause problems depending on how the EDID data is managed. Specifically, with several display devices connected to a splitter, which EDID information should be reported back to the graphics adapter? Low cost HDMI splitters have been found to exhibit unpredictable behavior. Typically they pass on EDID information from one of the attached devices, but exactly how the choice is determined is often unknown. Higher cost devices may have explicit schemes (for example they report the EDID from the display plugged into output port one), but these can cost more than the Rift itself! Generally, the use of third party splitters and video switching

boxes means that Rift EDID data may not be reliably reported to the OS, and therefore libovr will not be able to identify the Rift.

A.4.4 Windows: Direct3D enumeration

As described above the nature of the Rift display means that it is not suited to displaying the Windows desktop. As a result you might be inclined to set standalone mode for your regular monitor to remove the Rift from the list of devices displaying the Windows desktop. Unfortunately this also causes the device to no longer be enumerated when querying for output devices during Direct3D setup. As a result, the only viable option currently is to use the Rift in extended display mode.

B Chromatic Aberration

Chromatic aberration is a visual artifact seen when viewing images through lenses. The phenomenon causes colored fringes to be visible around objects, and is increasingly more apparent as our view shifts away from the center of the lens. The effect is due to the refractive index of the lens varying for different wavelengths of light (shorter wavelengths towards the blue end of the spectrum are refracted less than longer wavelengths towards the red end). Since the image being displayed on the Rift is composed of individual red, green, and blue pixels,³ it too is susceptible to the unwanted effects of chromatic aberration. The manifestation, when looking through the Rift, is that the red, green, and blue components of the image appear to be scaled out radially, and by differing amounts. Exactly how apparent the effect is depends on the image content and to what degree users are concentrating on the periphery of the image versus the center.

B.1 Correction

Fortunately, pixel shaders provide the means to significantly reduce the degree of visible chromatic aberration, albeit at some additional GPU expense. This is done by pre-transforming the image so that after the lens causes chromatic aberration, the end result is a more normal looking image. This is exactly analogous to the way in which we pre-distort the image to cancel out the distortion effects generated by the lens.

B.2 Sub-channel aberration

Although we can reduce the artifacts through the use of a pixel shader, it's important to note that this approach cannot remove them completely for an LCD display panel. This is due to the fact that each color channel is actually comprised of a range of visible wavelengths, each of which is refracted by a different amount when viewed through the lens. As a result, although we're able to distort the image for each channel to bring the peak frequencies back into spatial alignment, it's not possible to compensate for the aberration that occurs within a color channel. Typically, when designing optical systems, chromatic aberration across a wide range wavelengths is managed by carefully combining specific optical elements (in other texts, for example, look for “achromatic doublets”).

B.3 Shader implementation

The pixel shader for removing lens distortion was shown in Figure 7. This applies a radial distortion based on the distance of the display pixel from the center of distortion (the point at which the axis through the lens intersects the screen). Fortunately, chromatic aberration correction is also dependent on the same radial distance, and so we can achieve it with minimal modifications to the distortion shader.

The approach we employ applies an additional radial scale to the red and blue color channels in addition to the lens distortion. The scaling function for lens distortion was introduced in (13) and is repeated below:

$$f(r) = k_0 + k_1 r^2 + k_2 r^4 + k_3 r^6.$$

The additional red and blue scaling is expressed as follows:

³In a typical LCD display, pixel color is achieved using red, green, and blue dyes, which selectively absorb wavelengths from the back-light.

$$f'_R(r) = f(r)(c_0 + c_1 r^2) \quad (17)$$

$$f'_G(r) = f(r) \quad (18)$$

$$f'_B(r) = f(r)(c_2 + c_3 r^2). \quad (19)$$

The use of a quadratic term provides additional flexibility when tuning the function for medium and high values of r .

The following code shows the resulting pixel shader for correcting both lens distortion and chromatic aberration:

```
Texture2D Texture : register(t0);
SamplerState Linear : register(s0);
float2 LensCenter;
float2 ScreenCenter;
float2 Scale;
float2 ScaleIn;
float4 HmdWarpParam;
float4 ChromAbParam;

float4 main(in float4 oPosition : SV_Position, in float4 oColor : COLOR,
            in float2 oTexCoord : TEXCOORD0) : SV_Target
{
    float2 theta = (oTexCoord - LensCenter) * ScaleIn; // Scales to [-1, 1]
    float rSq = theta.x * theta.x + theta.y * theta.y;
    float2 thetal = theta * (HmdWarpParam.x + HmdWarpParam.y * rSq +
                            HmdWarpParam.z * rSq * rSq +
                            HmdWarpParam.w * rSq * rSq * rSq);

    // Detect whether blue texture coordinates are out of range
    // since these will scaled out the furthest.
    float2 thetaBlue = thetal * (ChromAbParam.z + ChromAbParam.w * rSq);
    float2 tcBlue = LensCenter + Scale * thetaBlue;

    if (any(clamp(tcBlue, ScreenCenter - float2(0.25, 0.5),
                  ScreenCenter + float2(0.25, 0.5)) - tcBlue))
        return 0;

    // Now do blue texture lookup.
    float blue = Texture.Sample(Linear, tcBlue).b;

    // Do green lookup (no scaling).
    float2 tcGreen = LensCenter + Scale * thetal;
    float green = Texture.Sample(Linear, tcGreen).g;

    // Do red scale and lookup.
    float2 thetaRed = thetal * (ChromAbParam.x + ChromAbParam.y * rSq);
    float2 tcRed = LensCenter + Scale * thetaRed;
    float red = Texture.Sample(Linear, tcRed).r;

    return float4(red, green, blue, 1);
}
```

Note the following:

- We apply the red and blue scaling after the lens distortion to reduce the amount of shader computation.
- We determine the texture coordinates for the blue component first because this extends farther than the other color channels and enables us to exit early if the texture coordinates lie outside of the range of the ‘half-screen’.

- The shader now has to perform three texture lookups instead of one, and only one color component from each lookup is used.